

Use natural language to create charts and graphs

Estimated time needed: **30** minutes

Overview

Imagine you are a data analyst or a data scientist of a marketing team at an e-commerce company. The company needs to understand customer purchasing behaviors over the last year to tailor their upcoming holiday campaigns. Traditionally, this would involve complex SQL queries, data wrangling in Python, and perhaps building visual dashboards to interpret the results including analyzing spreadsheets, creating charts, and maybe even some statistical analysis—tasks that require considerable time and expertise.

With the integration of Langchain and LLMs, you can simply ask, "Show me a visualization of monthly sales trends by product category," or "Generate a heatmap of customer activity by region." The system would use the `create_pandas_dataframe_agent` to process the CSV data, and then dynamically generate visualizations such as line graphs, bar charts, or heatmaps in response to these queries. This not only speeds up the data analysis process but also allows team members who may not be tech-savvy to engage directly with the data and make informed decisions quickly. This approach fosters a more collaborative environment and ensures that strategic decisions are backed by real-time data insights, visually represented for easy comprehension.

In this lab, you will learn how to seamlessly integrate data visualization into your conversational data analysis using Langchain and LLMs. Starting with CSV file data, you will use the `create_pandas_dataframe_agent` to build an interactive agent that not only understands and responds to your queries but also translates data responses into visual formats. You will explore how to dynamically generate charts, graphs, and heatmaps directly in response to natural language questions. This capability will enable you to visualize trends, compare figures, and spot patterns immediately, making your data analysis workflow both efficient and visually engaging. By the end of this project, you will have the skills to create a data conversational agent that acts as both analyst and visualizer, bringing data to life through dialogue.

In this lab we are going to use Llama 3 LLM hosted on the IBM watsonx.ai platform.

Table of contents

Exercises Exercise 1. Relationship between parental education level and student grades Exercise 2. Impact of internet access at home on grades Exercise 3. Explore LLM's code

Objectives

After completing the project, you should be able to:

- **Use Langchain with large language models:** Understand and apply the Langchain framework in conjunction with LLMs to interact with and analyze data stored in CSV files through natural language processing.
 - **Create conversational data agents:** Build a conversational agent that can understand and respond to natural language queries about data, enabling users to ask questions directly and receive immediate answers.
 - **Implement data visualization through dialogue:** Integrate data visualization tools within your conversational agent, allowing you to request and generate visual data representations such as graphs, charts, and heatmaps dynamically based on your queries.
 - **Enhance decision-making process:** Develop the capability to derive actionable insights from data via interactive dialogues and visual outputs, thereby improving the decision-making process and making data analysis accessible to non-technical stakeholders.
-

Setup

This project is based on Jupyter Notebook. If you're not familiar with it, here's a quick guide on how to run code within it:

A Jupyter Notebook consists of cells. To execute a code cell, click on the cell that you want to run and click the 'Run' button, as shown in the picture.

For this lab, we will be using the following libraries:

- `ibm-watson-ai` for using LLMs from IBM's watsonx.ai.
- `LangChain`, `langchain-ibm`, `langchain-experimental` for using its agent function to interact with data.
- `matplotlib` for additional plotting tools.
- `seaborn` for visualizing the data.

Installing required libraries

The following required libraries are **not** preinstalled in the Skills Network Labs environment. **You must run the following cell** to install them:

Note: We are pinning the version here to specify the version. It's recommended that you do this as well. Even if the library is updated in the future, the installed library could still support this lab work.

This might take approximately 1-2 minutes.

As we use `%%capture` to capture the installation, you won't see the output process. But after the installation completes, you will see a number beside the cell.

```
%%capture
!pip install ibm-watsonx-ai=="0.2.6"
!pip install langchain=="0.1.16"
!pip install langchain-ibm=="0.1.4"
!pip install langchain-experimental=="0.0.57"
!pip install matplotlib=="3.8.4"
!pip install seaborn=="0.13.2"
```

After you install the libraries, restart your kernel. You can do that by clicking the **Restart the kernel** icon.

Importing required libraries

We recommend you import all required libraries in one place (here):

```
# We use this section to suppress warnings generated by your code:
def warn(*args, **kwargs):
    pass
import warnings
warnings.warn = warn
warnings.filterwarnings('ignore')

from ibm_watsonx_ai.foundation_models import Model
from ibm_watsonx_ai.metanames import GenTextParamsMetaNames as
GenParams
from
ibm_watson_machine_learning.foundation_models.extensions.langchain
import WatsonxLLM

from langchain_experimental.agents.agent_toolkits import
create_pandas_dataframe_agent

import matplotlib.pyplot as plt
import pandas as pd
```

Data set

In this lab, you will work on the Student Alcohol Consumption data set `student-mat.csv` by UCI Machine Learning as an example. For more information, see [Kaggle](#). It is based on data collected from two secondary schools in Portugal. The students included in the survey were in mathematics and Portuguese courses.

The data set we are using is for the mathematics course. The number of mathematics students involved in the collection was 395. The data collected in locations such as Gabriel Pereira and Mousinho da Silveira includes several pertinence values. Examples of such data are records of demographic information, grades, and alcohol consumption.

Field	Description
school	GP/MS for the student's school
sex	M/F for gender
age	15-22 for the student's age
address	U/R for urban or rural, respectively
famsize	LE3/GT3 for less than or greater than three family members
Pstatus	T/A for living together or apart from parents, respectively
Medu	0 (none) / 1 (primary-4th grade) / 2 (5th - 9th grade) / 3 (secondary) / 4 (higher) for mother's education
Fedu	0 (none) / 1 (primary-4th grade) / 2 (5th - 9th grade) / 3 (secondary) / 4 (higher) for father's education
Mjob	'teacher,' 'health' care related, civil 'services,' 'at_home' or 'other' for the student's mother's job
Fjob	'teacher,' 'health' care related, civil 'services,' 'at_home' or 'other' for the student's father's job
reason	reason to choose this school (nominal: close to 'home', school 'reputation', 'course' preference or 'other')
guardian	mother/father/other as the student's guardian
traveltime	1 (<15mins) / 2 (15 - 30 mins) / 3 (30 mins - 1 hr) / 4 (>1hr) for a time from home to school
studytim	1 (<2hrs) / 2 (2 - 5hrs) / 3 (5 - 10hrs) / 4 (>10hrs) for weekly study time
failures	1-3/4 for the number of class failures (if more than three, then record 4)
schools	yes/no for extra educational support
sup	
famsup	yes/no for family educational support
paid	yes/no for extra paid classes for Math or Portuguese
activities	yes/no for extra-curricular activities
nursery	yes/no for whether attended nursery school
higher	yes/no for the desire to continue studies
internet	yes/no for internet access at home
romantic	yes/no for relationship status
famrel	1-5 scale on quality of family relationships
freetime	1-5 scale on how much free time after school
goout	1-5 scale on how much student goes out with friends
Dalc	1-5 scale on how much alcohol consumed on weekdays
Wal	
Wal	1-5 scale on how much alcohol consumed on the weekend
health	1-5 scale on health condition
absences	0-93 number of absences from school
G1	0-20 for the first-period grade

Field	Description
G2	0-20 for the second-period grade
G3	0-20 for the final grade

Load the data set

Execute the code in the following cell to load your dataset. This code reads the CSV file into a pandas DataFrame, making the data accessible for processing in Python.

```
df = pd.read_csv(
    "https://cf-courses-data.s3.us.cloud-object-
    storage.appdomain.cloud/ZNoKMJ9rssJn-QbJ49k0zA/student-mat.csv"
)
```

Let's examine the first five rows of the dataset to get a glimpse of the data structure and its contents.

```
df.head(5)
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob
0	GP	F	18	U	GT3	A	4	4	at_home
1	GP	F	17	U	GT3	T	1	1	at_home
2	GP	F	15	U	LE3	T	1	1	at_home
3	GP	F	15	U	GT3	T	4	2	health
4	GP	F	16	U	GT3	T	3	3	other

	famrel	freetime	goout	Dalc	Walc	health	absences	G1	G2	G3
0	4	3	4	1	1	3	6	5	6	6
1	5	3	3	1	1	3	4	5	5	6
2	4	3	2	2	3	3	10	7	8	10
3	3	2	2	1	1	5	2	15	14	15
4	4	3	2	1	2	5	4	6	10	10

[5 rows x 33 columns]

We can also review the detailed information for each column in the dataset, focusing on the presence of null values and the specific data types of each column.

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 395 entries, 0 to 394
Data columns (total 33 columns):
```

#	Column	Non-Null Count	Dtype
0	school	395 non-null	object
1	sex	395 non-null	object
2	age	395 non-null	int64
3	address	395 non-null	object
4	famsize	395 non-null	object
5	Pstatus	395 non-null	object
6	Medu	395 non-null	int64
7	Fedu	395 non-null	int64
8	Mjob	395 non-null	object
9	Fjob	395 non-null	object
10	reason	395 non-null	object
11	guardian	395 non-null	object
12	traveltime	395 non-null	int64
13	studytime	395 non-null	int64
14	failures	395 non-null	int64
15	schoolsup	395 non-null	object
16	famsup	395 non-null	object
17	paid	395 non-null	object
18	activities	395 non-null	object
19	nursery	395 non-null	object
20	higher	395 non-null	object
21	internet	395 non-null	object
22	romantic	395 non-null	object
23	famrel	395 non-null	int64
24	freetime	395 non-null	int64
25	goout	395 non-null	int64
26	Dalc	395 non-null	int64
27	Walc	395 non-null	int64
28	health	395 non-null	int64
29	absences	395 non-null	int64
30	G1	395 non-null	int64
31	G2	395 non-null	int64
32	G3	395 non-null	int64

dtypes: int64(16), object(17)

memory usage: 102.0+ KB

Load LLM

Execute the code in the cell below to load the llama-3-70b LLM model from watsonx.ai.

Additionally, we will configure the LLM to interact with data by integrating it with Langchain's `create_pandas_dataframe_agent`.

```
# Create a dictionary to store credential information
credentials = {
    "url"      : "https://us-south.ml.cloud.ibm.com"
}
```

```

# Indicate the model we would like to initialize. In this case, Llama
3 70B.
model_id      = 'meta-llama/llama-3-70b-instruct'

# Initialize some watsonx.ai model parameters
params = {
    GenParams.MAX_NEW_TOKENS: 256, # The maximum number of tokens
that the model can generate in a single run.
    GenParams.TEMPERATURE: 0,    # A parameter that controls the
randomness of the token generation. A lower value makes the generation
more deterministic, while a higher value introduces more randomness.
}
project_id    = "skills-network" # <--- NOTE: specify "skills-network"
as your project_id
space_id      = None
verify       = False

# Launch a watsonx.ai model
model = Model(
    model_id=model_id,
    credentials=credentials,
    params=params,
    project_id=project_id,
    space_id=space_id,
    verify=verify
)

# Integrate the watsonx.ai model with the langchain framework
llm = WatsonxLLM(model = model)

agent = create_pandas_dataframe_agent(
    llm,
    df,
    verbose=False,
    return_intermediate_steps=True # set
return_intermediate_steps=True so that model could return code that it
comes up with to generate the chart
)

```

Interact with your data

Let's start with a simple interaction.

Ask LLM how many rows of data are in the CSV file.

```

response = agent.invoke("how many rows of data are in this file?")
response['output']

```

```
'395'
```

From the output above, the model reports that there are 395 rows of data in the file.

Let's verify this count using Python code to ensure accuracy.

```
len(df)
```

```
395
```

The row count matches and is correct!

Curious about the code the LLM generated and used to create this result?

Run the code in the cell below to reveal the underlying commands.

```
response['intermediate_steps'][0][0].tool_input.replace('; ', '\n')
'len(df)'
```

Surprisingly, the LLM uses the same code as we do.

Also, we could let LLM return some data that we are looking for based on the CSV file.

```
response = agent.invoke("Give me all the data where student's age is over 18 years old.")
```

```
print(response)
```

```
{'input': "Give me all the data where student's age is over 18 years old.", 'output': 'The final answer is the dataframe with the students who are over 18 years old.', 'intermediate_steps':
[(AgentAction(tool='python_repl_ast', tool_input="df[df['age'] > 18]", log="Thought: I need to filter the dataframe to get the rows where the age is greater than 18.\nAction: python_repl_ast\nAction Input: df[df['age'] > 18]"),
school sex age address famsize Pstatus
Medu Fedu Mjob Fjob \
127 GP F 19 U GT3 T 0 1 at_home
other
153 GP M 19 U GT3 T 3 2 services
at_home
210 GP F 19 U GT3 T 3 3 other
other
247 GP M 22 U GT3 T 3 1 services
services
257 GP M 19 U LE3 A 4 3 services
at_home
270 GP F 19 U GT3 T 3 3 other
services
296 GP F 19 U GT3 T 4 4 health
```


other										
304	GP	M	19	U	GT3	T	3	3	other	
other										
306	GP	M	20	U	GT3	A	3	2	services	
other										
307	GP	M	19	U	GT3	T	4	4	teacher	
services										
308	GP	M	19	R	GT3	T	3	3	other	
services										
309	GP	F	19	U	LE3	T	1	1	at_home	
other										
310	GP	F	19	U	LE3	T	1	2	services	
services										
311	GP	F	19	U	GT3	T	2	1	at_home	
other										
312	GP	M	19	U	GT3	T	1	2	other	
services										
313	GP	F	19	U	LE3	T	3	2	services	
other										
314	GP	F	19	U	GT3	T	1	1	at_home	
health										
315	GP	F	19	R	GT3	T	2	3	other	
other										
336	GP	F	19	R	GT3	A	3	1	services	
at_home										
340	GP	F	19	U	GT3	T	2	1	services	
services										
350	MS	M	19	R	GT3	T	1	1	other	
services										
353	MS	M	19	R	GT3	T	1	1	other	
other										
370	MS	F	19	U	LE3	T	3	2	services	
services										
376	MS	F	20	U	GT3	T	4	2	health	
other										
383	MS	M	19	R	GT3	T	1	1	other	
services										
387	MS	F	19	R	GT3	T	2	3	services	
other										
390	MS	M	20	U	LE3	A	2	2	services	
services										
392	MS	M	21	R	GT3	T	1	1	other	
other										
394	MS	M	19	U	LE3	T	1	1	other	
at_home										
		reason	guardian	traveltime	studytime	failures	schoolsup			
famsup	\									
127	course	other		1	2	3	no			

yes							
153	home	mother	1	1	3	no	
yes							
210	reputation	other	1	4	0	no	
yes							
247	other	mother	1	1	3	no	
no							
257	reputation	mother	1	2	0	no	
yes							
270	home	other	1	2	2	no	
yes							
296	reputation	other	2	2	0	no	
yes							
304	home	other	1	2	1	no	
yes							
306	course	other	1	1	0	no	
no							
307	reputation	other	2	1	1	no	
yes							
308	reputation	father	1	2	1	no	
no							
309	reputation	other	1	2	1	yes	
yes							
310	home	other	1	2	1	no	
no							
311	other	other	3	2	0	no	
yes							
312	course	other	1	2	1	no	
no							
313	reputation	other	2	2	1	no	
yes							
314	home	other	1	3	2	no	
no							
315	reputation	other	1	3	1	no	
no							
336	home	other	1	3	1	no	
no							
340	home	other	1	3	1	no	
no							
350	home	other	3	2	3	no	
no							
353	home	other	3	1	1	no	
yes							
370	home	other	2	2	2	no	
no							
376	course	other	2	3	2	no	
yes							
383	other	mother	2	1	1	no	
no							

387	course	mother	1	3	1	no
no						
390	course	other	1	2	2	no
yes						
392	course	other	1	1	3	no
no						
394	course	father	1	1	0	no
no						
paid activities nursery higher internet romantic famrel freetime						
goout \						
127	no	no	no	no	no	3 4
2						
153	no	no	yes	no	yes	4 5
4						
210	yes	yes	yes	yes	no	4 3
3						
247	no	no	no	yes	yes	5 4
5						
257	no	no	yes	yes	no	4 3
1						
270	yes	yes	yes	yes	no	4 3
5						
296	yes	yes	yes	yes	no	2 3
4						
304	no	yes	yes	yes	yes	4 4
4						
306	no	yes	yes	yes	no	5 5
3						
307	yes	no	yes	yes	yes	4 3
4						
308	no	yes	yes	yes	no	4 5
3						
309	no	yes	no	yes	no	4 4
3						
310	no	yes	no	yes	no	4 2
4						
311	no	no	yes	no	yes	3 4
1						
312	no	no	no	yes	yes	4 5
2						
313	yes	no	no	yes	yes	4 2
2						
314	no	no	no	yes	yes	4 1
2						
315	no	no	yes	yes	yes	4 1
2						
336	yes	no	yes	yes	no	5 4
3						

340 4	yes		yes	yes	yes	yes	yes	4	3
350 4	no		no	yes	yes	yes	no	5	4
353 4	no		no	yes	yes	yes	no	4	4
370 2	no		yes	yes	yes	no	yes	3	2
376 3	yes		no	no	yes	yes	yes	5	4
383 2	no		no	yes	yes	no	no	4	3
387 2	no		yes	no	yes	yes	no	5	4
390 4	yes		no	yes	yes	no	no	5	5
392 3	no		no	no	yes	no	no	5	5
394 3	no		no	yes	yes	yes	no	3	2
	Dalc	Walc	health	absences	G1	G2	G3		
127	1	1	5	2	7	8	9		
153	1	1	4	0	5	0	0		
210	1	2	3	10	8	8	8		
247	5	5	1	16	6	8	8		
257	1	1	1	12	11	11	11		
270	3	3	5	15	9	9	9		
296	2	3	2	0	10	9	0		
304	1	1	3	20	15	14	13		
306	1	1	5	0	17	18	18		
307	1	1	4	38	8	9	8		
308	1	2	5	0	15	12	12		
309	1	3	3	18	12	10	10		
310	2	2	3	0	9	9	0		
311	1	1	2	20	14	12	13		
312	2	2	4	3	13	11	11		
313	1	2	1	22	13	10	11		
314	1	1	3	14	15	13	13		
315	1	1	3	40	13	11	11		
336	1	2	5	12	14	13	13		
340	1	3	3	4	11	12	11		
350	3	3	2	8	8	7	8		
353	3	3	5	4	8	8	8		
370	1	1	3	4	7	7	9		
376	1	1	3	4	15	14	15		
383	1	3	5	0	6	5	0		
387	1	2	5	0	7	5	0		
390	4	5	4	11	9	9	9		

392	3	3	3	3	10	8	7	
394	3	3	5	5	8	9	9)]]}

Let's get the code LLM used for charting this plot.

```
response['intermediate_steps'][0][0].tool_input.replace('; ', '\n')
"df[df['age'] > 18]"
```


Plot your data with natural language

Task 1

Generating a first visual on the data set to know the total number of male and female students in the data set.

We just tell the agent that "Plot the gender count with bars."

```
response = agent.invoke("Plot the gender count with bars.")
```



Let's see what code the LLM generated for plotting this chart.

```
print(response['intermediate_steps'][0][0].tool_input.replace('; ', '\n'))
```

```
df['sex'].value_counts().plot(kind='bar')
```

Task 2

Generating a pie chart for displaying the average value of weekend alcohol for each gender in the data set.

We will use the prompt "Generate a pie chart to display average value of Walc for each gender."

You may notice that the model generates two charts. The charts actually indicate the progressive improvement of the agent's code as it searches for the best way to answer to our prompt. It presents an improvement in the response to our query.

```
response = agent.invoke("Generate a pie chart to display average value  
of Walc for each Gender.")
```

Let's get the code LLM used for charting this plot.

```
print(response['intermediate_steps'][0][0].tool_input.replace('; ', '\n'))
```

Task 3

We can explore the impact of free time on grades based on the data.

```
response = agent.invoke("Create box plots to analyze the relationship  
between 'freetime' (amount of free time) and 'G3' (final grade) across  
different levels of free time.")
```

Execute the code below to retrieve the Python script the LLM used for plotting.

```
print(response['intermediate_steps'][0][0].tool_input.replace('; ', '\n'))
```

Task 4

We can explore the effect of alcohol consumption on academic performance.

```
response = agent.invoke("Generate scatter plots to examine the  
correlation between 'Dalc' (daily alcohol consumption) and 'G3', and  
between 'Walc' (weekend alcohol consumption) and 'G3'.")
```

Execute the code below to retrieve the Python script the LLM used for plotting.

```
print(response['intermediate_steps'][0][0].tool_input.replace('; ', '\n'))
```

Exercises

Exercise 1 - Relationship between parental education level and student grades

```
# your code here
```

Exercise 2 - Impact of internet access at home on grades

```
# your code here
```

Exercise 3 - Explore LLM's code

Can you find what code the model used to generate the plot for exploring the relationship between absences and academic performance?

You could run the corresponding code and from the response chain, you could see the code used from charting.

```
# your code here
```

```
# the code that model use here
```

Authors

[Kang Wang](#)

Kang Wang is a Data Scientist in IBM. He is also a PhD Candidate in the University of Waterloo.

[Wojciech Fulmyk](#) Wojciech "Victor" Fulmyk is a Data Scientist at IBM. He is also a PhD Candidate in Economics in the University of Calgary.

Other contributors

[Ricky Shi](#) Ricky Shi is a data scientist at the Ecosystems Skills Network at IBM.

<!--## Change Log--!>

| <!-- | Date (YYYY-MM-DD) | Version | Changed By |
|------------|-------------------|-----------------------------|------------------------------|
| 2024-05-10 | 0.2 | Kang Wang & Wojciech Fulmyk | Initial version created |
| 2024-02-23 | 0.1 | Elio Di Nino | Update library documentation |

Copyright © 2024 IBM Corporation. All rights reserved.